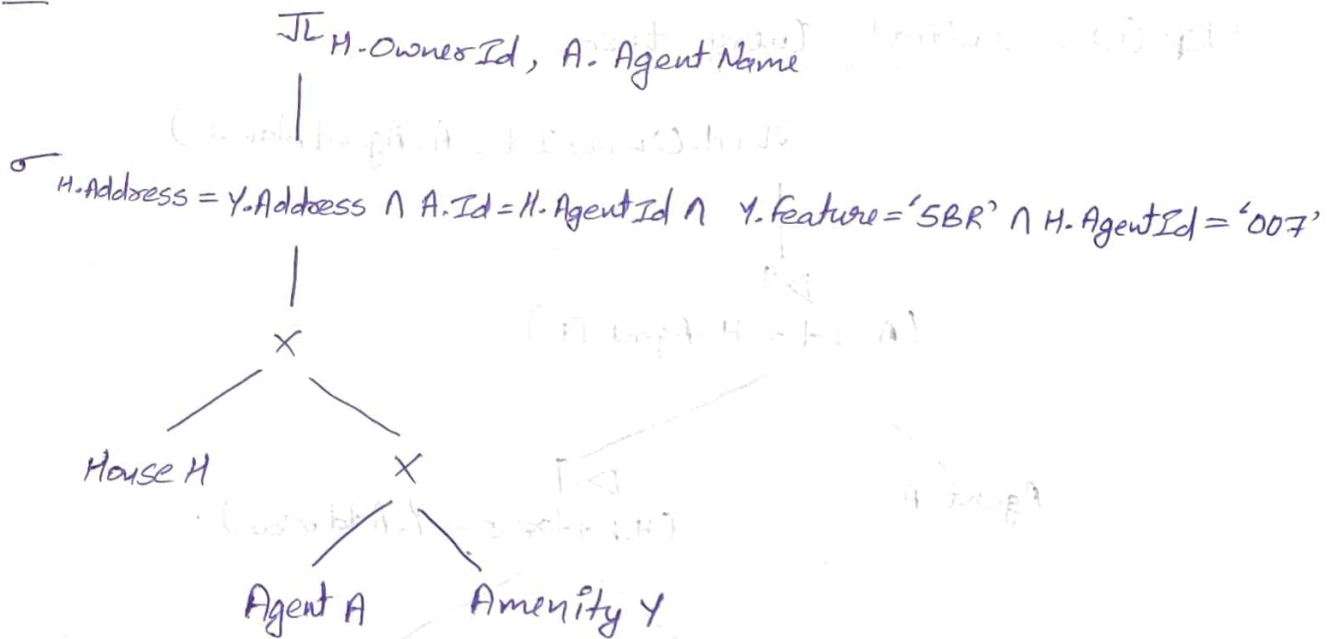


Answer 1:

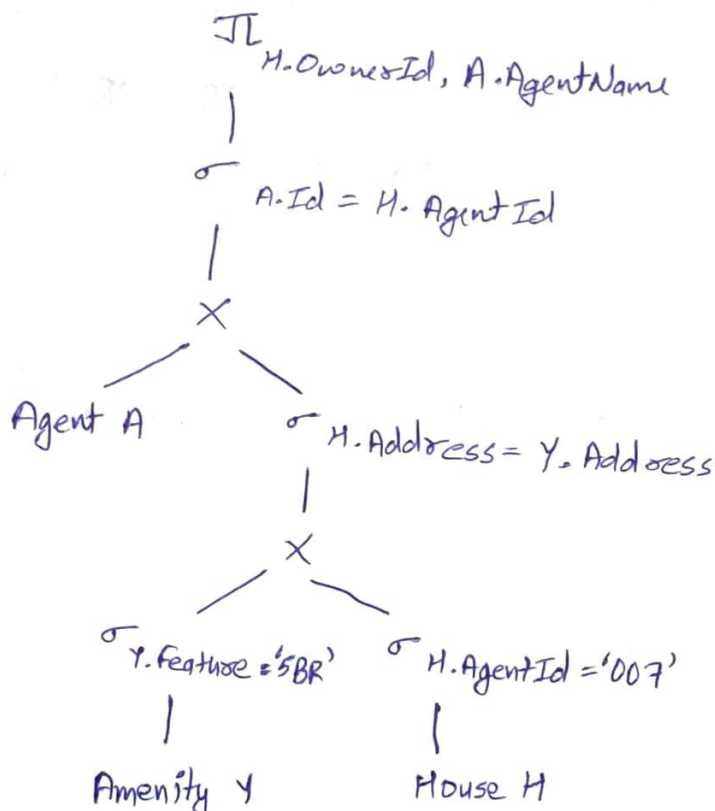
SELECT H.OwnerId, A.AgentName FROM House H, Agent A, Amenity Y
WHERE H.Address = Y.Address AND A.Id = H.AgentId AND
Y.Feature = 'SBR' AND H.AgentId = '007';

$\pi_{H.OwnerId, A.AgentName} (\sigma_{H.Address = Y.Address \wedge A.Id = H.AgentId \wedge Y.Feature = 'SBR' \wedge H.AgentId = '007'} (House H \times Agent A \times Amenity Y))$

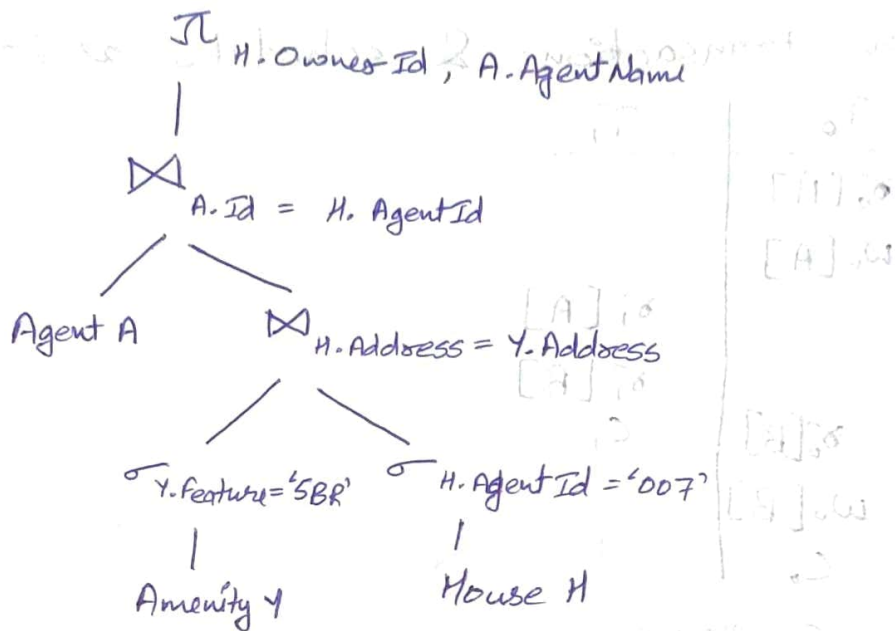
Step 1



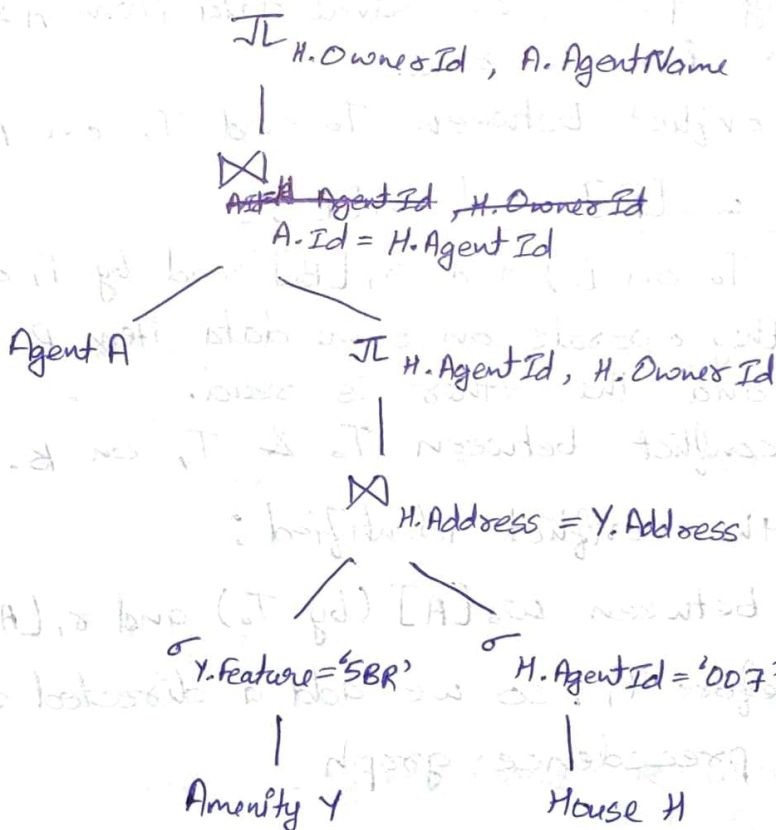
Step 2 : Moving Select operation down, and rearrange leaf



Step 3: Combine Cartesian Product and select into join



Step 4: Moving Projection Down



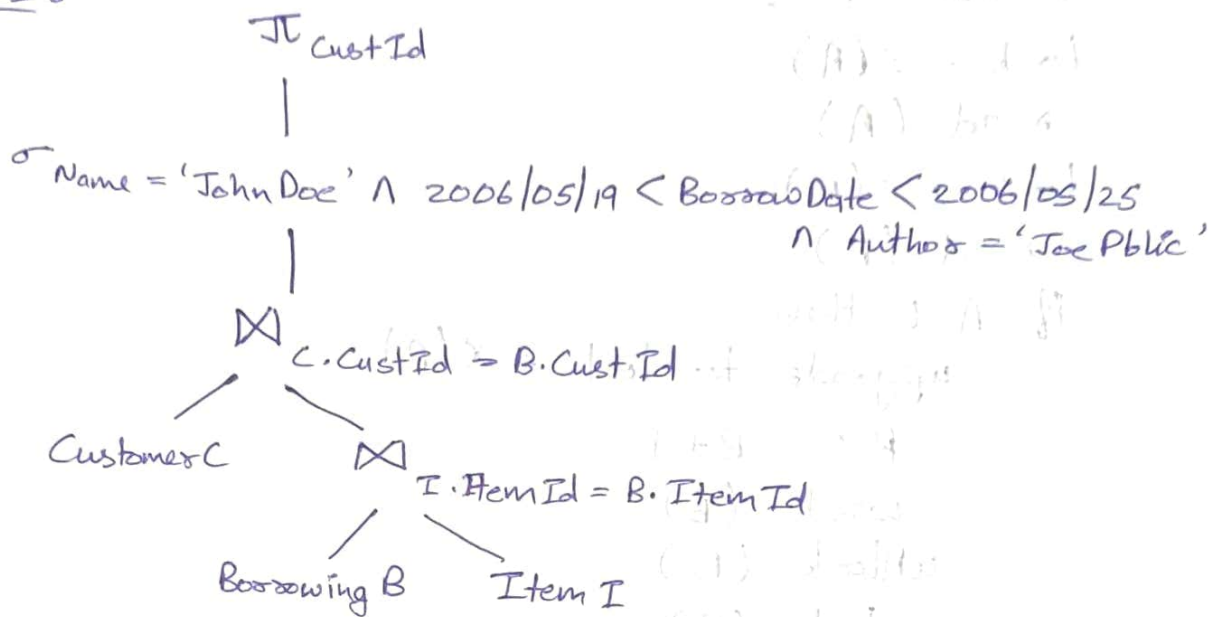
Final Expression -

$$\pi_{H.OwnerId, A.AgentName} \left(\text{Agent A} \bowtie_{A.Id = H.AgentId} \left(\left(\sigma_{Y.Feature = 'SBR'}(Amenity Y) \right) \bowtie_{H.Address = Y.Address} \left(\sigma_{H.AgentId = '007'}(House H) \right) \right) \right)$$

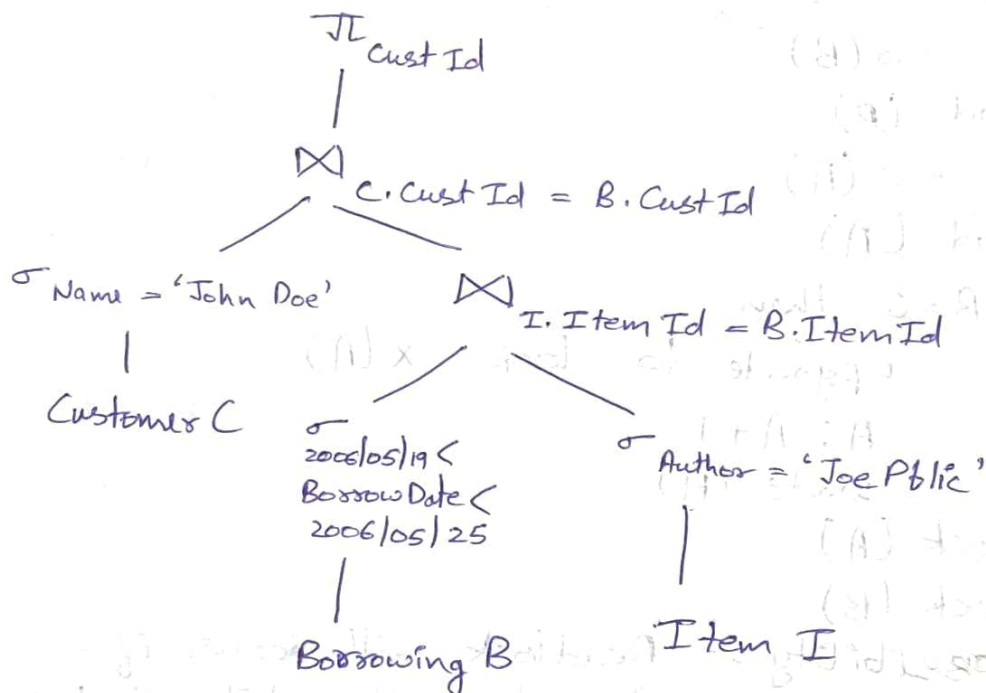
Answer 2:

$\pi_{CustId}(\sigma_{Name = 'JohnDoe' \wedge 2006/05/19 < BorrowDate < 2006/05/25 \wedge Author = 'JoePblc'}(Customer \bowtie Borrowings \bowtie Item))$

Step 1:

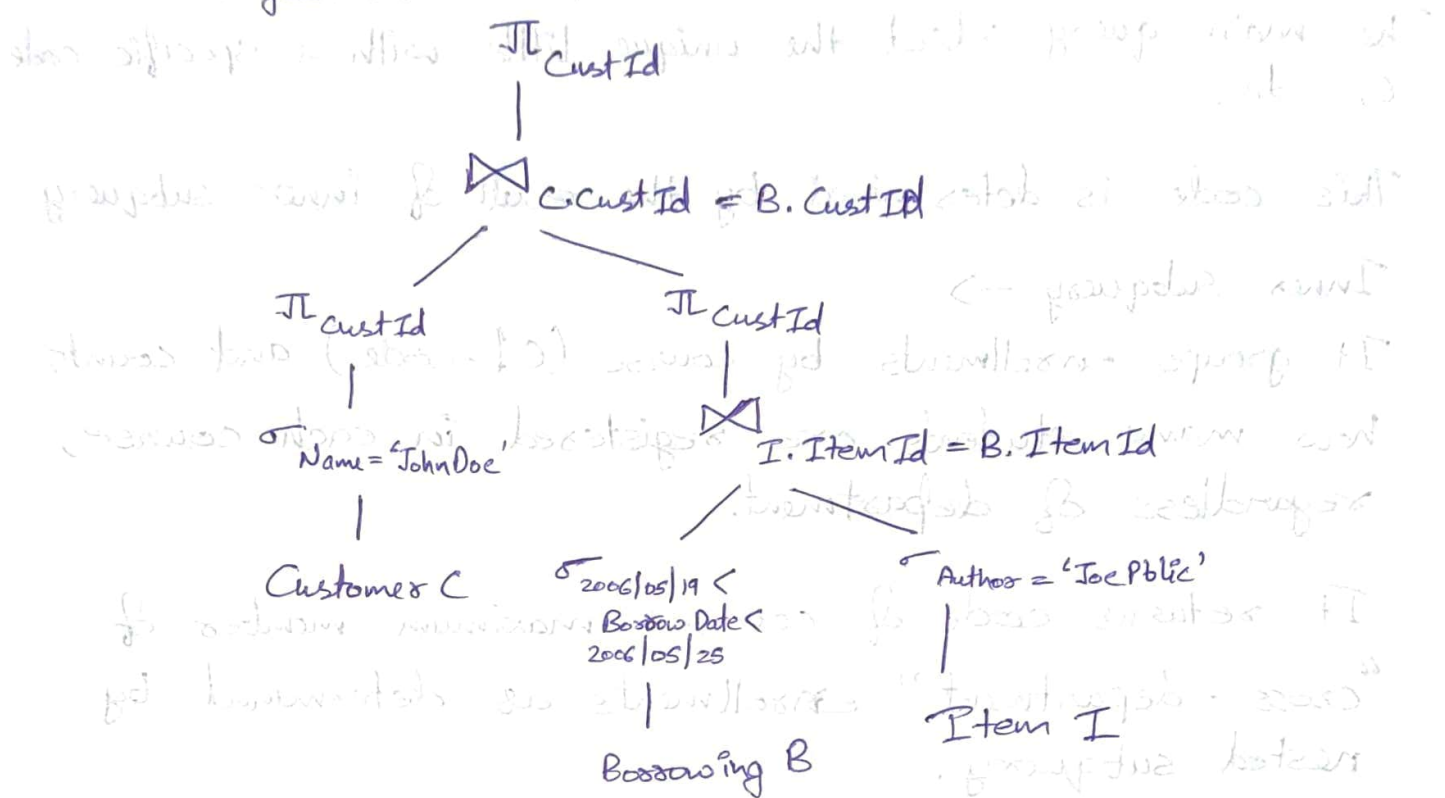


Step 2: Move Select down



Step 3:-

Move Projection down



Relational algebra expression for the query:

$$\pi_{C.CustId} \left(\left(\left(\pi_{C.CustId} \left(\sigma_{Name = 'JohnDoe'} (Customer\ C) \right) \right) \bowtie C.CustId = B.CustId \right. \right. \\ \left. \left(\pi_{B.CustId} \left(\left(\sigma_{2006/05/19 < BorrowDate < 2006/05/25} (Borrowing\ B) \right) \bowtie I.ItemId = B.ItemId \right. \right. \right. \\ \left. \left. \left(\sigma_{Author = 'Joe Public'} (Item\ I) \right) \right) \right) \right)$$

SQL

```
SELECT C.CustId FROM (
    (SELECT C.CustId FROM Customer C WHERE C.Name = 'John Doe')
    JOIN
    (SELECT B.CustId
     FROM
        (SELECT * FROM Borrowing B WHERE BorrowDate > '2006/05/19'
         AND BorrowDate < '2006/05/25')
        JOIN
        (SELECT * FROM Item I WHERE Author = 'Joe Public')
        ON I.ItemId = B.ItemId
    )
    ON C.CustId = B.CustId ) ;
```


Answer 3:

The main query selects the unique title with a specific code C_code.

This code is determined by the result of inner subquery

Inner Subquery →

It groups enrollments by course (C1_code) and counts how many students are registered in each course, regardless of department.

It returns code of courses maximum member of "cross-department" enrollments as determined by nested subquery.

Nested Subquery →

It calculates the maximum count of cross-department enrollments for any course, means that the course (being taken) by students of different department than the one offering the course (C2_dept <> S2_dept)

Final →

The query finally retrieves the title of the course that has the maximum number of students from departments other than offering the course.

Answer 4:

Given two transactions & schedules as :-

T_0	T_1
$r_0[A]$	
$w_0[A]$	
	$r_1[A]$
	$r_1[B]$
$r_0[B]$	c_1
$w_0[B]$	
c_0	

Identifying Conflict:

Conflict between $w_0[A]$ and $r_1[A]$

$w_0[A]$ (write by T_0 on A) and $r_1[A]$ (read by T_1 on A) conflicts because they operate on same data item A & one is write & another one is read.

\therefore we have a conflict between T_0 and T_1 on A .

Conflict between $w_0[B]$ and $r_1[B]$:

$w_0[B]$ (write by T_0 on B) and $r_1[B]$ (read by T_1 on B) conflict because they operate on same data item B and one is a write and the other is read.

\therefore we have a conflict between T_0 & T_1 on B .

Now, based on the conflict identified:

① from the conflict between $w_0[A]$ (by T_0) and $r_1[A]$ (by T_1): T_0 must come before T_1 , so we add a directed edge from $T_0 \rightarrow T_1$ in the precedence graph.

② from the conflict between $w_0[B]$ (by T_0) and $r_1[B]$ (by T_1): T_0 must come before T_1 , so we add another directed edge from $T_0 \rightarrow T_1$. Therefore the precedence graph will look as:-
 $T_0 \rightarrow T_1$

Now, Since precedence graph does not contains any cycles, so the given schedule is conflict serializable.

Answer 5:

Transaction T_1 :

lock - S(A)

read (A)

lock - S(B)

read (B)

if $A=0$ then

upgrade to lock - X(B)

$B := B + 1$

write (B)

unlock (B)

unlock (A)

Transaction T_2 :

lock - S(B)

read (B)

lock - S(A)

read (A)

if $B=0$ then

upgrade to lock - X(A)

$A := A + 1$

write (A)

unlock (A)

unlock (B)

Deadlock Possibility: - Deadlock will occur if -

T_1 locks A and waits for lock B, while T_2 locks B and waits to lock A.

This is called circular wait and is a classic deadlock condition.

If both transaction acquire their initial locks but then wait on each other's locked resources, neither can proceed, resulting in a deadlock.